

CS1510 Dynamic Programming Problems 22 & 27, Reduction Problems 3 & 4

Rebecca Negley, Sean Myers

October 17, 2011

22. The input to this problem is n points x_1, \dots, x_n on a line. A good path P has the property that one endpoint of P is the origin and every x_i is covered by P . Note that P need not be simple. That is, it can backtrack over territory that it has already covered. Assume a vehicle moves along this path from the origin at unit speed. The response time r_i for each x_i is the time until the vehicle first reaches x_i . The problem is to find the good path that minimizes $\sum_{i=1}^n r_i/n$, the average response time. Give a polynomial time algorithm for this problem.

Solution: Suppose at time t , a path P is at point x_p . The next unvisited point that P visits must be the unvisited point closest to x_p on the left or the unvisited point closest to x_p on the right. If P instead goes to some point farther away, it must pass through one of those points on the way, so one of those two points will in fact be the next one visited. Create two sequences x_{1_1}, \dots, x_{1_k} and x_{2_1}, \dots, x_{2_m} ($k+m=n$) such that $x_{1_i} \geq 0$ for all $i \in [1, k]$ and $x_{2_j} < 0$ for all $j \in [1, m]$. Additionally, $x_{1_i} \leq x_{1_{i+1}}$ for all $i \in [1, k-1]$ and $x_{2_j} \geq x_{2_{j+1}}$ for all $j \in [1, m-1]$. Then, since the path starts at the origin (between the two sequences), the closest unvisited point to the left of the current point is $\min_j(x_{2_j} | x_{2_j} \text{ unvisited})$. Similarly, the closest unvisited point to the right of the current point is $\min_i(x_{1_i} | x_{1_i} \text{ unvisited})$.

Now we can construct our tree. At level lvl , each node corresponds to a sequence of lvl points. Each node has two children, one corresponding to the the smallest unvisited x_{1_i} (corresponding to $\min i$) being added to the sequence and one corresponding to the greatest unvisited x_{2_j} (corresponding to $\min j$) being added to the sequence. Then, we have a binary tree with n levels, so we have 2^n possible solutions at the leaves. To reduce complexity, we will prune the tree as follows:

- (a) If there are two nodes on the same level that have sequences ending at the same point and have the same closest unvisited points on the left and the same closest unvisited points on the right, prune the one with the greater average response time.

With this pruning rule, each level can have at most $(k+1) * (m+1) < n^2$ possible pairs of closest unvisited right point and closest unvisited left point. We can create our tree in an $n \times n \times n \times 2 \times 2$ matrix MGP , the first dimensions corresponding to the level of the tree, the second to the index of the closest unvisited point to the right, third to the index of the closest unvisited point to the left, the fourth to whether the path is at the right or left input (0 for right, 1 for left), and the fifth to whether we are storing the minimum average response time or the current response time. We initialize each location $MGP[\alpha, \beta, \gamma, 0/1, 0]$ to $+\infty$ except for $MGP[0, 1, 1, 0, 0]$, $MGP[0, 1, 1, 0, 1]$, $MGP[0, 1, 1, 1, 0]$, and $MGP[0, 1, 1, 1, 1]$, which we set to 0. Then, we traverse down the levels of our matrix (or tree) by building each level from the previous one. If $MGP[\alpha, \beta, \gamma, 0, 0]$ has a value $< \infty$, then we can set $MGP[\alpha+1, \beta, \gamma+1, 0, 0]$ to the minimum of its current value and $(\alpha * MGP[\alpha, \beta, \gamma, 0, 0] + (MGP[\alpha, \beta, \gamma, 0, 1] + (x_{1_{\gamma+1}} - x_{1_\gamma}))) / (\alpha + 1)$. If we change this location from its current value, set $MGP[\alpha+1, \beta, \gamma+1, 0, 1]$ to $MGP[\alpha, \beta, \gamma, 0, 1] + (x_{1_{\gamma+1}} - x_{1_\gamma})$, the response time at that point. Similarly, we can set the location

coordinating to a left traversal. Then, when we are finished filling in our table, we seek return the minimum of $\text{MGP}[n, k+1, m+1, 0, 0]$ and $\text{MGP}[n, k+1, m+1, 1, 0]$, corresponding to the minimum of the path using all the points and ending at the right endpoint and the path using all the points and ending at the left endpoint. In code, our algorithm looks as follows:

```

minGoodPath( $x_1, \dots, x_n$ )
  construct  $x_{1_1}, \dots, x_{1_k}$  ordered (L-R) sequence of nonnegative points
  construct  $x_{2_1}, \dots, x_{2_m}$  ordered (R-L) sequence of negative points
  construct MGP of size  $n \times n \times n \times 2 \times 2$ 
  initialize MGP to  $+\infty$ 
  MGP[0,1,1,0,0]=0
  MGP[0,1,1,0,1]=0
  MGP[0,1,1,1,0]=0
  MGP[0,1,1,1,1]=0
  for  $lvl$  from 1 to  $n$  do:
    for  $i$  from 1 to  $k+1$  do:
      for  $j$  from 1 to  $m+1$  do:
        if MGP[lvl,  $i, j, 0, 0$ ]  $< \infty$  do:
          MGP[lvl+1,  $i+1, j, 0, 0$ ] = min{MGP[lvl+1,  $i+1, j, 0, 0$ ], ( $lvl * \text{MGP}[lvl, i, j, 0, 0]$ 
            + MGP[lvl,  $i, j, 0, 1$ ] + ( $x_{1_i} - x_{1_{i-1}}$ )) / ( $lvl + 1$ )}
          if MGP[lvl+1,  $i+1, j, 0, 0$ ] changed:
            MGP[lvl+1,  $i+1, j, 0, 1$ ] =  $lvl * \text{MGP}[lvl, i, j, 0, 0]$  + MGP[lvl,  $i, j, 0, 1$ ] + ( $x_{1_i} - x_{1_{i-1}}$ )
            MGP[lvl+1,  $i, j+1, 1, 0$ ] = min{MGP[lvl+1,  $i, j+1, 1, 0$ ], ( $lvl * \text{MGP}[lvl, i, j, 0, 0]$ 
              + MGP[lvl,  $i, j, 0, 1$ ] + ( $x_{1_{i-1}} - x_{2_j}$ )) / ( $lvl + 1$ )}
            if MGP[lvl+1,  $i, j+1, 1, 0$ ] changed:
              MGP[lvl+1,  $i, j+1, 1, 1$ ] =  $lvl * \text{MGP}[lvl, i, j, 0, 0]$  + MGP[lvl,  $i, j, 0, 1$ ] + ( $x_{1_{i-1}} - x_{2_j}$ )
        if MGP[lvl,  $i, j, 1, 0$ ]  $< \infty$  do:
          MGP[lvl+1,  $i+1, j, 0, 0$ ] = min{MGP[lvl+1,  $i+1, j, 0, 0$ ], ( $lvl * \text{MGP}[lvl, i, j, 1, 0]$ 
            + MGP[lvl,  $i, j, 1, 1$ ] + ( $x_{1_i} - x_{2_{j-1}}$ )) / ( $lvl + 1$ )}
          if MGP[lvl+1,  $i+1, j, 0, 0$ ] changed:
            MGP[lvl+1,  $i+1, j, 0, 1$ ] =  $lvl * \text{MGP}[lvl, i, j, 1, 0]$  + MGP[lvl,  $i, j, 1, 1$ ] + ( $x_{1_i} - x_{2_{j-1}}$ )
            MGP[lvl+1,  $i, j+1, 1, 0$ ] = min{MGP[lvl+1,  $i, j+1, 1, 0$ ], ( $lvl * \text{MGP}[lvl, i, j, 1, 0]$ 
              + MGP[lvl,  $i, j, 1, 1$ ] + ( $x_{2_{j-1}} - x_{2_j}$ )) / ( $lvl + 1$ )}
            if MGP[lvl+1,  $i, j+1, 1, 0$ ] changed:
              MGP[lvl+1,  $i, j+1, 1, 1$ ] =  $lvl * \text{MGP}[lvl, i, j, 1, 0]$  + MGP[lvl,  $i, j, 1, 1$ ] + ( $x_{2_{j-1}} - x_{2_j}$ )
  return min(MGP[ $n, k+1, m+1, 1, 0$ ], MGP[ $n, k+1, m+1, 0, 0$ ])

```

27. Give a polynomial time algorithm for the following problem. The input consists of a two dimensional array R of non-negative integers, and an integer k . The value $R_{t,p}$ gives the number of users requesting page p at time t . At each integer time, the server can broadcast either 0 or 1 pages. If the server broadcasts page p at time t , the requests of all the users who requested page p strictly before time t are satisfied. The server can make at most k broadcasts. The goal is to pick the k times to broadcast and the pages to broadcast at those k times in order to minimize the total time (over all requests) that requests/users have to wait in order to have their requests satisfied.

Solution: We are going to set up the tree in the following way: Each level represents some time t , and there are n levels. For each node, there are a lot of possible branches. Each branch stands for broadcasting a certain page at the next time interval, so there are $P+1$ branches (p pages, and not broadcasting anything). This will produce a tree of $(p+1)^n$ leaves, where the leaf of weight k , and the minimum of all those k weights, is the optimal solution.

The first pruning rule, which was already mentioned is to prune any tree with weight greater than k .

There can only be k broadcasts, hence any more pages added would be an infeasible solution. The second pruning rule is if two nodes have the exact same nodes in them (on the same level), but have different user waiting requests. Choose the node with the least request waiting time. The reasoning for this second rule is like so: Let's say we have two nodes: a and b where $a < b$ in terms of user requests pending. Let's say b has the optimal solution, such that $time(b) + time(j) + ...time(n)$ is the optimal time. Then you could add the same times to a , such that: $time(a) + time(j) + ...time(n) < time(b) + time(j) + ... + time(n)$, which means a is now the most optimal. Therefore b can be pruned.

With both these pruning rules, we see that there can be a combination of p pages times k groups that those pages can be in (duplicates allowed). This means there will be $k * p$ leaves.

Now let's pseudocode this! First we need an array that can hold the $p * k$ possibilities, at each level, so it will be the array `request[time, pages, groups]`. This array will be initialized to infinity, except for `request[0,0,0]`, which will be set to 0. After that, the next step is to fill the tree. To do that, we must traverse the times, starting at 0, and using the layer above to fill out the one below. At each level above, we must transfer all the times of the current broadcasts, which would look like: `request[lvl+1, k, p] = min(request[lvl+1, k, p], request[lvl, k, p] + cost(lvl+1, p))`. The min is so we don't wipe out the best value if it is already there. The cost is a separate function that will add all the requests up to that point since we are not broadcasting anything (which means it will cost all the current requests still active) Another thing we need to do is calculate if we add a certain page to a group at that time: `request[lvl+1, k, p] = min(request[lvl+1, k, p], request[lvl, k, p] + cost(lvl+1, p) - request[lvl+1, p])`. what this is saying is the same as above except that we subtract the requests of the page from the cost (since we chose the page as a broadcast and not as a stagnant page).

The final algorithm looks something like this:

```
Initialize request[time, groups, pages] = infinity
request[0,0,0] = 0

for(lvl = 0 to time):
    for(k = 1 to groups):
        for(p = 1 to pages):
            request[lvl+1, k, p] = min(request[lvl+1, k, p], request[lvl, k, p] + cost(lvl+1, p))
            request[lvl+1, k, p] = min(request[lvl+1, k, p], request[lvl, k, p] + cost(lvl+1, p) - request[lvl+1, p])

cost(level, page):
    for(0 to all pages):
        add requests of all pages, multiply it by level.
    return sum
```

3. Show that if there is an $O(n^k)$, $k \geq 1$, times algorithm for squaring a degree n polynomial, then there is an $O(n^k)$ time algorithm for multiplying two degree n polynomials. Assume that the polynomials are given by their coefficients.

Solution: The first thing we must do is set up the single polynomial equation. If there are two polynomials a and b , then the equations should be arranged $a + b$. Setting this up takes $2n$ time, just to go through the n elements of each. Running this through the squaring algorithm takes $4n^k$ time. Finally, we must extrapolate the results. If they are multiplied correctly, such that the first coefficient is multiplied to all other coefficients of itself, then the second, and so on, then to extrapolate the data would be to look at the position of the multiplication. For example, let's say $a = x_1 + x_2$ and $b = x_3 + x_4$ where the x 's are coefficients of some polynomials and we want $a * b$. We would set it up so: $c = x_1 + x_2 + x_3 + x_4$, and then multiply it, by itself. To find what we want, we number

the positions (like they already are), and extrapolate based on that. So if we want $x_1 * x_3$, it is going to be the (first position-1)*2n + second position, which will give you the index at which the answer is at. So in this case, it will be (1-1)+3, or the 3rd multiplication.(which makes sense since the first multiplications would be: $x_1 * x_1 + x_1 * x_2 + x_1 * x_3$). This should also roughly take n time, which means the algorithm will run in $O(n^k)$ time, based on the squaring degree n polynomial.

4. Consider the following variant of the minimum Steiner tree problem. The input is n points in the plane. Each point is given by its Cartesian coordinates. The problem is to build a collection of roads between these points so that you can reach any city from any other city and the total length of the roads is minimized. The collection of roads should be output as an adjacency list structure. Show by reduction that if you can solve this problem in linear time, then you can sort n numbers in linear time.

Solution: We will reduce the sorting problem to the minimum Steiner tree problem variant. For each input number x in the sorting problem, we create a point $(x, 0)$ to input into the Steiner tree problem. To do this for each x takes linear time. Then, we input these n points into the variant Steiner tree problem and receive an adjacency list back in linear time (by supposition). Then, we look through the adjacencies of the n input points (linear time) to see which two points have only one adjacent point. These must be the endpoints, so we take the point with the lesser x value as the left endpoint. Then, we construct the rest of our sorted list of numbers by following our adjacency list (which also takes linear time). In code, our reduction looks like this:

```
Sort( $x_1, \dots, x_n$ )
  for  $i$  from 1 to  $n$  do:
     $p_i = (x_i, 0)$  //  $O(n)$ 
  adjList = variantSteiner( $p_1, \dots, p_n$ ) //  $O(n)$ 
   $j_1, j_2$  = points in adjList with only one adjacency
   $x_{\alpha(1)} = \min(\text{x val of } j_1, \text{x val of } j_2)$ 
  for  $i$  from 2 to  $n$  do:
     $x_{\alpha(i)} = x_j$  such that  $(x_j, 0)$  adjacent to  $(x_{i-1}, 0)$  and  $x_j > x_{i-1}$  //  $O(n)$ 
  return  $x_{\alpha(1)}, \dots, x_{\alpha(n)}$ 
```

Hence, if a linear time algorithm exists for the variant Steiner problem, a linear algorithm exists for sorting numbers.